

DAY 2 GLOSSARY

Topic models, and how a BERT model reads text.

Quasi-sentence

The Manifesto Project's unit of coding. One statement, one code.

cmp_code

The numeric Manifesto category, such as 504 for Welfare State Expansion.

RILE

The standard left-right scale. 13 codes on each side.

Preprocessing

Lowercasing, dropping stopwords, and counting, before any model runs.

Stopwords

Words carrying no topic signal. We also drop country and party names.

Document-term matrix

Rows are documents, columns are words, cells are counts.

CountVectorizer

The sklearn tool that builds that matrix.

TF-IDF

Counts, downweighted for words appearing everywhere. Day 3's baseline.

LDA

Latent Dirichlet Allocation. Finds groups of words that travel together.

Topic

A distribution over words. A policy area, and not a left-right score.

STM

Structural Topic Model. Lets a covariate like year shape topic prevalence.

rpy2

The bridge letting us call the R stm package from inside Python.

Tokenizer

Splits text into the chunks the model reads, and maps them to id numbers.

Token

One of those chunks. Often a whole word, sometimes a word piece.

input_ids

The list of id numbers for your sentence.

attention_mask

1 for a real token, 0 for padding. Tells the model what to ignore.

Padding

Adding filler so every row in a batch is the same length.

Truncation

Cutting text that runs past the limit.

512 tokens

The wall for BERT-family models. Longer documents get chunked.

[CLS] and [SEP]

Markers the tokenizer adds at the start and between segments.

Classification head

The small layer bolted on top that turns numbers into labels.

Logits

The raw scores the head produces, before any normalising.

Softmax

Turns logits into probabilities summing to 1.

Argmax

The position of the largest value, which is the predicted label.