

REPRODUCIBILITY CHECKLIST

Everything to record before you write up a fine-tuned measure.

A fine-tuned model is a **measurement instrument**, so it needs the same documentation any other instrument gets. Record these before you write up, because most of them are impossible to reconstruct afterwards.

The data

- Source and version. For us, Manifesto Project main dataset **MPDS2026a** and corpus **2026-1**.
- Which countries, elections and parties you selected, and the rule you used to select them.
- How many documents and how many quasi-sentences you ended with.
- The label scheme. For RILE, the 13 codes per side and the v5 subcategory rule that sends 202.2, 605.2 and 703.2 to uncoded.
- Whether you may redistribute the data. For the Manifesto corpus you may not, so ship code.

The split

- Train and test sizes, per class.
- The seed used for sampling.
- Confirmation that no training sentence appears in the test set.

The model

- Backbone id, exactly as written, such as **sentence-transformers/all-MiniLM-L6-v2**.
- Method and library version, such as SetFit, and the transformers version.
- Batch size, epochs, and any other argument you changed from the default.
- The training seed.
- Hardware, such as a Colab T4, and roughly how long training took.

The evidence

- Held-out Cohen's kappa and macro F1 for the fine-tune.
- The same two numbers for your TF-IDF baseline, on the same test set.
- The confusion matrix, so a reader can see which class it misses.
- Where the model is saved, and its Hub id if you pushed it.
- A model card naming the task, the data, the score and the seed.

The one-line test

Hand this page and your notebook to a colleague with their own Manifesto key. If they cannot land within a point or two of your kappa, something on this list is missing.